

Why forgetting is the *price of remembering cheaply.*

Honest version: Transformers are that friend who saves every screenshot. Great memory, dead phone storage. Recurrent models are the friend who keeps one whiteboard and wipes it. We kept arguing about which friend BDH is. Spoiler: it picked the whiteboard — on purpose. Come break it with us.

OUR BET • PROVE US WRONG

KV-cache: one drawer per word. Grows forever, never lies. **Whiteboard: N slots, total.** Cheap forever, quietly lies when two words want the same corner. If you can keep the whiteboard honest at 500 words in 16 slots, we lose.

Built by Cloud 9, DataForge. If you know attention a little, you're good. 6 mins. The toy below is ours — scrappy, live, definitely not the real BDH.

HOW THIS PAGE WORKS

No lecture. Just mess with it.

1 • STARE

It's already running — 120 pairs, 32 slots. Cabinet's chilling at 100%. Whiteboard's already messed up ~80 times. Click a yellow cell. Ask it what happened.

2 • BE MEAN

Press "Flood it". That's 500 words into 16 slots — hostel-wifi-energy. Then flip the **state rule**: hash vs mamba-gate vs delta vs hola vs bdh-p. Same flood, five different ways to die.

3 • OH, SO BDH DID THIS TOO

Yeah. The real machine chose the whiteboard life. Same headache, fancier math. Scroll and we'll show you the receipts.

THE LAB • LIVE TOY

Two memories, one stream of garbage words.

We feed both the same junk — mango-3 → 42, chai-7 → 88, you get it. Cabinet files everything. Whiteboard goes $\text{hash}(\text{key}) \% N$ and whoever wrote last wins — like a mess register where the last entry counts. Pick a key below and catch the whiteboard lying.

PAIRS STREAMING IN • N



120 pairs

WHITEBOARD SLOTS • N

32 ▾

STATE RULE • 5 WAYS TO PAY

INTERROGATE KEY

hash — last wins

▾


lotus-0

▾

hash: cartoon of raw overwrite. Last writer wins the slot. Honest baseline — real attention interferes by superposition, not modulo, same shape. Cartoons of write/read rules for teaching — not real kernels. Toy ≠ official models.

Filing cabinet — KV-cache $O(n)$ • exact

Slots held: 120



Grows one per token. Retrieval: dot-product over everything. Cost: you pay for every old token, forever.

Whiteboard — recurrent state 0(1) · lossy

Slots: 32 · overwrites: 88



Yellow = someone got erased here. Red = currently lying to you. Click a cell to cross-examine.

Truth beside estimate

Key	Truth	KV says	Whiteboard says	Who stole it?
lotus-0 drawer 9 · hash	65	65	84 ✕	overwritten by orbit-5, rickshaw-3, monsoon-2


Rule hash recall: 26.7% · Cabinet: 100% · Slot collisions: 88

THE BILL · LIVE CURVES

Exactness sends an invoice.

Left: drawers you have to buy. Right: answers you get wrong. Both drawn live from your N — and yeah, they print fine now, we learned that the hard way (canvas, never again).


Slots vs tokens



Right now: cabinet 120 drawers vs whiteboard 32 under rule hash. At 500: 500 vs 32.

Whiteboard accuracy vs stream length

accuracy per rule – live (hash)



Rule hash: $n=120$, $N=32 \rightarrow 24.1\%$. Switch rules at same n/N — that gap is the whole talk.

BDH MODULE • THE PART WE ACTUALLY SWEATED OVER

Okay, so the real machine did the whiteboard thing too.

Who: Dragon Hatchling (BDH), Pathway's weird-and-cool Post-Transformer. Quick note so nobody gets yelled at in Q&A: **not** a Mamba-style SSM — BDH-GPU is its own GPU-friendly rewrite from ReLU low-rank maps + linear attention. BDH doesn't keep drawers; it scribbles into **synapses while reading**. The one equation we're confident enough to frame (Eq.5):

$\rho_{t-1,l} = \sum_{\tau < t} v_{\tau,l-1}^* x_{\tau,l}^T U^{t-\tau}$ — BDH-GPU attention state, Eq.5: fixed-shape matrix, additive updates, $O(1)$ in length

Same deal as your whiteboard upstairs, just with learned keys, a decay U , and ~5% of neurons awake at a time. If you want the full math: Kosowski et al. §3.2 Eqs.4–5, Hebbian lens 6–8, toy code App.E. We read it so you can skim it.

Receipt #1 • vendor told us, we didn't verify: 134M BDH gets **95% at 32K** $\rightarrow$ **82% at 128K** on BABILong QA1–QA5 (pathway.com). Holds long, sags under load — exactly your toy's curve, at real scale. Pending contamination / independent / leaderboard review. A benchmark is not a deployment, we'll say it in the viva too.

Receipt #2 · the cousin BDH-CQ: 150M, **29.5% pass@2 on ARC-AGI-1 @ ~\$0.0007/task** (arXiv:2608.09888). Thinks in latent space, no chain-of-thought essays — cheap, not champion. Self-reported, weights unreleased.

Don't leave with the wrong lesson (we almost did): this forgetting is silent — no "sorry, forgot" flag. And remembering for one chat ≠ learning forever. Getting fast memory into slow weights is still unsolved. Also please, never call BDH an SSM in front of the judges.

LANDSCAPE IN 30 SECONDS

Who pays what?

Family	Memory	Recall	Bill	Why care
Transformer KV	$O(n)$	perfect	$O(T^2)$, OOMs	gold standard, expensive
Mamba SSM '24	$O(1)$ gated	lossy far	cheap	throughput king, weak needle
HOLA '26	$O(1)$ +cache	recovered	cheap+cache	hippocampus for surprises
BDH / CQ	$O(1)$ synaptic	lossy, sags	lowest \$/task shown	memory = reasoning fabric

60-SECOND CHECK · BE HONEST

Bet you can't call the next move.

No marks, no tracking. Judges will basically ask you this in the viva, so practice getting it wrong here first.

- ☐ **n=500, N=16:** whiteboard stays perfect because hashing is uniform.
- ☐ **n=500, N=16:** whiteboard collapses; too many keys share too few slots.
- ☐ **KV-cache:** forgets oldest tokens first to stay $O(1)$.

Answer, then set N=16 and flood it to check yourself.

RECEIPTS • VERIFY, DON'T TRUST

Sources & caps.

1. Kosowski et al. *Dragon Hatchling: Missing Link between Transformer and Brain*. arXiv:2509.26507 (30 Sep 2025). Eqs.4–8, BDH-GPU App.E. [arxiv](#)
2. Engdahl, Kosowski et al. *BDH-CQ: In-Context Learning with Recurrent Latent Reasoning*. arXiv:2608.09888 (10 Aug 2026). $S_t = U(S_{t-1}, D_t)$. [arxiv](#)
3. Cui. *Hippocampus for Linear Attention*. arXiv:2607.02303 (2 Jul 2026). $O(1)$ but lossy; overwrites. [arxiv](#)
4. Gu & Dao. *Mamba*. 2024. SSM baseline — related, distinct from BDH.

Caps, honestly: toy = string-hash % N, int values, exact-match. No learned keys, decay, sparsity. Small-N curves are live; BABILong/ARC are precomputed vendor numbers. Toy $\neq$ official BDH. Code: team-written with AI drafting help (Muse Spark) for code, prose, and design; team traced, tested, and defends every line.

Built by Cloud 9 for DataForge — if something looks off, it's probably our hash function, tell us in Q&A. Single file, no login, no secrets · [repo](#) · [live](#)